CS211 Fall 2022 Programming Assignment I

David Menendez

Due: Wednesday, October 12, 2022

This assignment is designed to give you some initial experience with programming in C, as well as compiling, running, and debugging. Your task is to write seven small C programs.

Section 1 describes the six programs, section 2 describes how your project will be graded, and section 3 describes how to structure and submit your project. In particular, section 3.1 describes how to set up your project and provides an introduction to using the auto-grader. Please read the entire assignment description before beginning the assignment.

Note that the assignment is due at 11:59 PM (ET). Submissions received after 11:59 PM will not receive full credit. You are strongly encouraged not to work until the last minute. Plan to submit your assignment by October 11.

Advice Pay attention to the program descriptions, especially how the programs receive their input. Some programs use only their command-line arguments (argv), others read from standard input (with scanf() or similar), or from a file. Read the specifications carefully, and think about what is being asked.

Make a plan before you start coding. If you cannot solve the problem yourself, you will not be able to write a program to do it. Don't be afraid to draw diagrams to guide your understanding of the data structures or control flow.

If you have not used command-line shells before, be aware that the text you type at the prompt is processed before being sent to a program as arguments. Most punctuation marks, aside from ., _, and -, have special meanings and must be escaped or enclosed with quotation marks in order to pass them to a program verbatim. You are not responsible for interpreting what the user typed: assume that any strings your program receives are intentional and perform only the processing required by the program specification.

Finally, it is not necessary to print a complete line at a time. If you are printing several strings or other values sequentially, it is both simpler and more efficient to call printf multiple times than to allocate a string, write the output to the string, and then print the string all at once. Similarly, it is easier to read input one token at a time rather than reading a line and the breaking the line into tokens.

1 Program descriptions

You will write six programs for this project. Except where explicitly noted, your programs may assume that their inputs are properly formatted. However, your programs should be robust. Your

program should not assume that it has received the proper number of arguments, for example, but should check and report an error where appropriate.

Except where noted, programs should always terminate with exit code EXIT_SUCCESS (that is, return 0 from main).

1.1 rot13 13: String operations

Rot-13 (rotate-13) is a simple substitution cipher, in which each letter is encoded by the letter appearing 13 places later (or earlier) in the alphabet. Thus, A becomes N, M becomes Z, N becomes A, and Z becomes M. Because 13 is half the number of letters in the alphabet, we use the same substitution for ciphering and deciphering. That is, applying rot-13 twice to a text yields the original text.

Write a program rot13 that encodes its argument using rot-13 and prints the results, followed by a newline character. You MAY assume that rot13 will receive exactly one argument, but it is good practice to check whether the number of arguments is correct. Only scenarios with one argument may be tested.

Note that rot-13 only affects alphabetic characters. All non-alphabetic characters (punctuation, whitespace, etc.) MUST be printed unchanged. Note also that arguments may be empty strings.

Usage

```
$ ./rot13 Hello
Uryyb
$ ./rot13 Uryyb
Hello
$ ./rot13 ebg13
rot13
$ ./rot13 ""

$ ./rot13 "Awesome Power!"
Njrfbzr Cbjre!
~/D/R/2/2/p/b/rot13 $ ./rot13 erzrzore\ gung\ gur\ furyy\ cebprffrf\ lbhe\ fgevat\ svefg
remember that the shell processes your string first
```

Notes You may find the functions in ctype.h helpful, especially isalpha().

rot13 can be written without using malloc. Recall that output can be printed a single character at a time using putchar(), so it is not necessary to print the entire output at once.

1.2 sorta: String operations II

Write a program sorta that sorts and prints its arguments. sorta takes zero or more arguments, sorts them lexicographically, and then prints each argument on its own line.

If sorta receives no arguments, it should exit without printing anything.

You are free to use any sorting algorithm you find convenient, but sorta MUST use strcmp() for string comparison. Note that strcmp() uses lexicographical ordering (so a string comes after its prefix) based on ASCII comparison of characters (so 'Z' comes before 'a'). (See man strcmp for usage information.)

Treat the arguments (excluding argv[0]) as given: make no attempt to process or otherwise normalize the strings before sorting. Note that your shell will perform some manipulation on your input before passing the strings to your program, and most punctuation marks have special meanings: use escape sequences and/or quotation marks to avoid this.

Usage

```
$ ./sorta foo bar baz quuz
bar
baz
foo
quux
$ ./sorta aaa aAa AaA AAA
AAA
AaA
aAa
aaa
$ ./sorta 3 20 100
100
20
3
$ ./sorta " z" a
 7.
```

1.3 list: Linked lists

Write a program list that maintains and manipulates a sorted linked list according to instructions received from standard input. The linked list is maintained in order, meaning that the items in the list are stored in increasing numeric order after every operation.

Note that list will need to allocate space for new nodes as they are created, using malloc; any allocated space should be deallocated using free as soon as it is no longer needed.

Note also that the list will not contain duplicate values.

list supports two operations:

insert n Adds an integer n to the list. If n is already present in the list, it does nothing. The instruction format is an i followed by a space and an integer n.

delete n Removes an integer n from the list. If n is not present in the list, it does nothing. The instruction format is a d followed by a space and an integer n.

After each command, list will print the length of the list followed by the contents of the list, in order from first (least) to last (greatest).

list must halt once it reaches the end of standard input.

Input format Each line of the input contains an instruction. Each line begins with a letter (either "i" or "d"), followed by a space, and then an integer. A line beginning with "i" indicates that the

integer should be inserted into the list. A line beginning with "d" indicates that the integer should be deleted from the list.

Your program will not be tested with invalid input. You may choose to have list terminate in response to invalid input.

Output format After performing each instruction, list will print a single line of text containing the length of the list, a colon, and the elements of the list in order, all separated by spaces.

Usage Because list reads from standard input, you may test it by entering inputs line by line from the terminal.

```
$ ./list
i 5
1 : 5
d 3
1 : 5
i 3
2 : 3 5
i 500
3 : 3 5 500
d 5
2 : 3 500
^D
```

To terminate your session, type Control-D at the beginning of the line. (This is indicated here by the sequence ^D.) This closes the input stream to list, as though it had reached the end of a file.

Alternatively, you may use input redirection to send the contents of a file to list. For example, assume list_commands.txt contains this text:

```
i 10
i 11
i 9
d 11
```

Then we may send this file to list as its input like so:

```
$ ./list < list_commands.txt
1 : 10
2 : 10 11
3 : 9 10 11
2 : 9 10</pre>
```

1.4 mexp: Matrix manipulation

Write a program mexp that multiplies a square matrix by itself a specified number of times. mexp takes a single argument, which is the path to a file containing a square $(k \times k)$ matrix M and a non-negative exponent n. It computes M^n and prints the result.

Note that the size of the matrix is not known statically. You must use malloc to allocate space for the matrix once you obtain its size from the input file.

To compute M^n , it is sufficient to multiply M by itself n-1 times. That is, $M^3 = M \times M \times M$. Naturally, a different strategy is needed for M^0 .

Input format The first line of the input file contains an integer k. This indicates the size of the matrix M, which has k rows and k columns.

The next k lines in the input file contain k integers. These indicate the content of M. Each line corresponds to a row, beginning with the first (top) row.

The final line contains an integer n. This indicates the number of times M will be multiplied by itself. n is guaranteed to be non-negative, but it may be 0.

For example, an input file file.txt containing

indicates that mexp must compute

$$\left[\begin{array}{ccc} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{array}\right]^2.$$

Output format The output of mexp is the computed matrix M^n . Each row of M^n is printed on a separate line, beginning with the first (top) row. The items within a row are separated by spaces. Using file.txt from above,

```
$ ./mexp file1.txt
30 36 42
66 81 96
102 126 150
```

1.5 bst: Binary search trees

Write a program bst that manipulates binary search trees. It will receive commands from standard input, and print resposes to those commands to standard output.

A binary search tree is a binary tree that stores integer values in its interior nodes. The value for a particular node is greater than every value stored its left sub-tree and less than every value stored in its right sub-tree. The tree will not contain any value more than once. bst will need to allocate space for new nodes as they are created using malloc; any allocated space should be deallocated using free before bst terminates.

This portion of the assignment has two parts.

Part 1 In this part, you will implement bst with three commands:

insert n Adds a value to the tree, if not already present. The new node will always be added as the child of an existing node, or as the root. No existing node will change or move as as result of inserting an item. If n was not present, and hence has been inserted, bst will print inserted. Otherwise, it will print not inserted. The instruction format is an i followed by a decimal integer n.

search n Searches the tree for a value n. If n is present, bst will print present. Otherwise, it will print absent. The instruction format is an s followed by a space and an integer n.

print Prints the current tree structure, using the format in section 1.5.1.

Part 2 In this part, you will implement bst with an additional fourth command:

delete n Removes a value from the tree. See section 1.5.2 for further discussion of deleting nodes. If n is not present, print absent. Otherwise, print deleted. The instruction format is a d followed by a space and an integer n.

Input format The input will be a series of lines, each beginning with a command character (i, s, p, or d), possibly followed by a decimal integer. When the input ends, the program should terminate.
Your program will not be tested with invalid input. A line that cannot be interpreted may be treated as the end of the input.

Output format The output will be a series of lines, each in response to an input command. Most commands will respond with a word, aside from **p**. The format for printing is described in section 1.5.1.

Usage

```
$ ./bst
i 1
inserted
i 2
inserted
i 1
not inserted
s 3
absent
p
(1(2))
```

As with list, the ^D here indicates typing Control-D at the start of a line in order to signal the end of file.

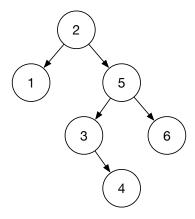


Figure 1: A binary search tree containing six nodes

1.5.1 Printing nodes

An empty tree (that is, NULL) is printed as an empty string. A node is printed as a (, followed by the left sub-tree, the item for that node, the right subtree, and), without spaces.

For example, the output corresponding to fig. 1 is ((1)2((3(4))5(6))).

1.5.2 Deleting nodes

There are several strategies for deleting nodes in a binary tree. If a node has no children, it can simply be removed. That is, the pointer to it can be changed to a NULL pointer. Figure 2a shows the result of deleting 4 from the tree in fig. 1.

If a node has one child, it can be replaced by that child. Figure 2b shows the result of deleting 3 from the tree in fig. 1. Note that node 4 is now the child of node 5.

If a node has two children, its value will be changed to the maximum element in its left subtree. The node which previously contained that value will then be deleted. Figure 2c shows the result of deleting 5 from the tree in fig. 1. Note that the node that previously held 5 has been relabeled 4, and that the previous node 4 has been deleted.

Note that the value being deleted may be on the root node.

1.6 rpn: A simple calculator

Write a program that performs integer calculations specified in Reverse Polish notation (RPN), a mathematical notation in which operators are given *after* the arguments to which they apply. Your program will need to maintain a stack of stored values for the operators to use as inputs, and will need to detect certain problems such as (1) the stack containing more than one value when the computation is complete and (2) the stack not containing enough values to perform the requested operation.

Input rpn takes a single argument, which will be a calculation written in the RPN format described in section 1.6.1.

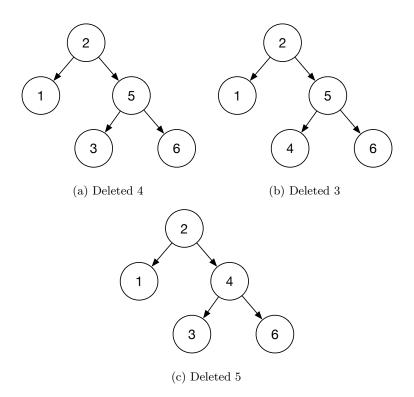


Figure 2: The result of deleting different values from the tree in fig. 1

Output On success, rpn prints the result of the computation and terminates with exit status EXIT_SUCCESS.

On failure, rpn terminates with exit status EXIT_FAILURE. For debugging purposes, it is recommended that rpn print a helpful error message, but our tests will only check the exit status. See section 1.6.2 for the precise error conditions rpn must detect.

Usage

```
$ ./rpn 20,30+
50
$ ./rpn 3,4,5x-
-17
$ ./rpn 3,4-5x
```

In the following examples, the specific wording of the error message is not required.

```
$ ./rpn 2,3++
```

Error: Too many operators

\$./rpn 2,3,4+

Error: Not enough operators

1.6.1 Notation Format

To specify a simple operation, such as 12 + 3, we write 12,3+. The input consists of two or more numbers, each followed by an operator symbol.

The comma (,) pushes the previous number onto the stack.

The math operators (+ - x /) pop the top two values off the stack; perform additon, subtraction, multiplication, or division; and push the answer onto the stack. A math operator following a number will push the number onto the stack before popping; a math operator following another operator will not push.

Note that we use x (the letter 'x') for multiplication, as * has a special meaning to the shell. You may additionally accept * for multiply, but this will not be tested.

We consider the stack to grow from left to right, so the value at the top of the stack becomes the right operand, and the value below it becomes the left operand. If a math operator follows a number For example, the string 10,4-2/ will be interpreted as:

- 1. Push 10 onto the stack
- 2. Push 4 onto the stack
- 3. Pop 4 and 10, push 6 onto the stack
- 4. Push 2 onto the stack
- 5. Pop 2 and 6, push 3 onto the stack

For the purposes of this assignment, you may assume that a comma will only occur between numbers.

1.6.2 Error conditions

In successful operation, there will always be at least two numbers in the stack when performing arithmetic and the stack will contain one number (the answer) after the last computation.

rpn MUST detect and report the following error conditions:

- 1. The stack does not contain two or more numbers when performing a math operator
- 2. The stack does not contain exactly one number after the final operator
- 3. The input does not end with an operator
- 4. The specified calculation divides by zero

When any of these error conditions are detected, rpn SHOULD print a helpful error message of your choosing, and MUST terminate with exit status EXIT_FAILURE.

There are additional error conditions which **rpn** MAY detect, but which will not be tested by the auto-grader, such as input containing characters other than digits or the five operators.

1.6.3 Checking the exit status

The auto-grader automatically checks your program's exit status. When running rpn yourself, you can confirm the exit status in a few ways.

First, most shells will set the environment variable \$? to the exit status of the last command. Use echo \$? to check its value after running rpn.

```
$ ./rpn 2,3x4,5x+
26
$ echo $?
0
```

Second, you can use && or || to perform a second command if the first one succeeds or fails, respectively.

```
$ ./rpn 17,2/ && echo success
8
success
$ ./rpn 17,2+- || echo failure
Non-specific error message
failure
```

1.6.4 Implementation notes

You may use any convenient data type to represent the stack, such as a linked list or array. If using an array, your program MUST NOT use a fixed size for the stack: you will need to use malloc() to allocate the array and either expand the array as needed or determine based on the input the bounds of the array. ¹

While you may use sscanf() to read the integers from the input string, it is fairly straightforward to read the integer yourself. Note that a character c contains a digit character (such as '5'), you can convert it to an integer by subtracting the ASCII value of '0' (that is, c - '0').

¹Since the array only expands after a comma, the maximum possible stack size is one more than the number of commas in the input.

```
PA1 Auto-grader, Release 1
```

```
Tests performed: 76 of 76
```

Tests failed: 0

Regular credit

	Points	${\tt Failed}$	${\tt Score}$
rot13	5.0	-	5.0
sorta	5.0	-	5.0
list	20.0	-	20.0
mexp	20.0	-	20.0
bst:1	15.0	-	15.0
bst:2	15.0	-	15.0
rpn:1	15.0	-	15.0
rpn:2	5.0	-	5.0
	100.0		100.0

Figure 3: Perfect submission

2 Grading

Your submission will be awarded up to 100 points, based on how many test cases your programs complete successfully.

Figure 3 shows the output from the full auto-grader for a perfect submission.

The auto-grader provided for students includes half of the test cases that will be used during grading. Thus, it will award up to 50 points (half the number of points in each category).

Make sure that your programs meet the specifications given, even if no test case explicitly checks it. It is advisable to perform additional tests of your own devising.

2.1 Academic integrity

You must submit your own work. You should not copy or even see code for this project written by anyone else, nor should you look at code written for other classes. We will be using state of the art plagiarism detectors. Projects which the detectors deem similar will be reported to the Office of Student Conduct.

Do not post your code on-line or anywhere publically readable. If another student copies your code and submits it, both of you will be reported.

3 Submission

Your solution to the assignment will be submitted through Sakai. You will submit a Tar archive file containing the source code and makefiles for your project. Your archive should not include any compiled code or object files.

The remainder of this section describes the directory structure (section 3.2), the requirements for your makefiles (section 3.3), how to create the archive (section 3.4), and how to use the provided auto-grader (section 3.5).

3.1 Getting started

Download the auto-grader and use tar to extract it. (The \$ in these examples indicates a prompt. The command you should type comes after the prompt and does not include the \$.)

```
$ tar -xf pa1-grader.tar
```

This will create a directory pa1/, containing the auto-grader script and its associated data files. First, create a subdirectory pa1/src/, which will contain a subdirectory for each of the six programs (see section 3.2 for the suggested layout).

```
$ cd pa1
pa1$ mkdir src
```

(Here we are changing the prompt to indicate the working directory. As before, you only type the text after the \$.)

Source directory and Makefile For each program, you will create a subdirectory with the same name as the program. It will contain a makefile and the source code for your project. For example, to start work on rot13, one could create a directory rot13 inside pa1/src/ and copy the makefile template into it.

```
pa1$ mkdir src/rot13
pa1$ cp template.make src/rot13/Makefile
```

Open the Makefile in the editor of your choice and replace program_name with the name of the program (for example, rot13).

Alternatively, you can make the copy and edit the program name in a single set using sed.

```
pa1$ sed '{s/program_name/rot13/;}' template.make > src/sorta/Makefile
```

Source file Once your makefile is created, create a file for your source code. The template makefile assumes your code will be a single file with the same name as the program. That is, the source for rot13 will be a file named rot13.c in the subdirectory src/rot13. You are permitted to modify the makefile to use multiple source files or different source file names, but take care to ensure compatibility with the auto-grader.

Build directory Before you start writing your program, make sure the auto-grader can compile and run your code. Put a minimal program in your source code file (for example, one that prints "Hello, world!" and terminates), and then run the auto-grader.

```
pa1$ ./grader.py --quiet
```

Since we expect all the tests to fail, we can specify quiet mode to minimize output. Check for any errors relating to make or the compiler being unable to find files.

Running the auto-grader will create a build directory, if it does not already exist, with sub-directories for each program. The auto-grader will compile your programs in these directories, not in the src directories. Doing so keeps your source directories free of clutter. When running your own tests, use the compiled program in the appropriate build directory and recompile using make.

You may use the --init option to create the build directories without compiling or running tests.

```
pa1$ ./grader.py --init
pa1$ cd build/rot13
pa1/build/rot13$ make
```

Because the build directory is created from your source code, you are always free to delete it and have the auto-grader reconstruct it.

3.1.1 Testing early and often

The auto-grader provided to you is the same one we will use to test your code. To avoid disaster, make sure that the auto-grader can compile and execute your program! Each time you make progress with a program, run the auto-grader to check for improvement or regression.

In the early stages of testing for a program, you can save some time by only running tests on that particular program and stopping after the first failed test.

```
pa1$ ./grader.py --stop list
```

The --stop or -1 option instructs the auto-grader to halt after the first failed test and to print your program's complete output. Specifying a program or test group instructs the auto-grader to skip tests for other programs.

Once you are confident that your program is mostly correct, you can omit --stop to see all the failed tests. To see the complete output from your program in failed test cases, use --verbose or -v. (See section 3.5 for additional auto-grader features.)

3.2 Directory structure

Your project should be stored in a directory named src, which will contain six sub-directories. Each subdirectory will have the name of a particular program, and contain (1) a makefile, and (2) any source files needed to compile your program. Typically, you will provide a single C file named for the program. That is, the source code for the program bst would be a file bst.c, located in the directory src/bst.

Figure 4 shows the layout of a typical project.

3.3 Makefiles

We will use make to manage compilation. Each program directory will contain a file named Makefile that describes at least two targets. The first target must compile the program. An additional target, clean, must delete any files created when compiling the program (typically just the compiled program).

```
src
+- bst
   +- Makefile
   +- bst.c
 +- list
     +- Makefile
    +- list.c
+- mexp
     +- Makefile
     +- mexp.c
 +- rot13
     +- Makefile
     +- rot13.c
 +- rpn
     +- Makefile
    +- rpn.c
 +- sorta
    +- Makefile
     +- sorta.c
```

Figure 4: Typical project layout

The auto-grader script is distributed with an example makefile, shown in fig. 5 (note that an actual makefile must use tabs rather than spaces for indentation).

It is simplest to copy this file into the directories for each program, replacing rot13 with the name of that specific program. This will ensure that you programs will be compiled with the recommended options.

It is further recommended that you use make to compile your programs, rather than invoking the compiler directly. This will ensure that your personal testing is performed with the same compiler settings as the auto-grader. The makefiles created in the build directory by the auto-grader refer to the makefiles you create in the source directory and therefore pick up any changes made.

You may add additional compiler options as you see fit, but you are advised to leave the compiler

Figure 5: Example makefile

warnings, sanitizers, and debugger information (-g). The makefile shown here specifies the C99 standard, in order to allow C++-style // comments; you may change that to C89, if you prefer.

Compiler options The sample makefile uses the following compiler options, listed in the CFLAGS make variable:

- -g Include debugger information, used by GDB and AddressSanitizer.
- -std=c99 Require conformance with the 1999 C Standard. (Disable GCC extensions.) You may change this to -std=c89 or -std=c90 at you discretion.
- -Wall Display most common warning messages.
- -Wvla Warn when using variable-length arrays.
- **-Werror** Promote all warnings to errors.
- -fsanitize=address,undefined Include run-time checks provided by AddressSanitizer and UBSan. This will add code that detects many memory errors and guards against undefined behavior. (Note that these checks discover problems with your code. Disabling them will not make your code correct, even if it seems to execute correctly.)

Target and dependency variables Note the use of \$@ (indicating the target name) and \$^ (indicating the dependencies). The auto-grader uses some advanced features of make to put the source files and object files in different directories. If you prefer to write your own Makefile instead of using the sample, you will need to use these variables in order for the auto-grader to successfully compile your project. Contact me with any questions about how to do this.

3.4 Creating the archive

We will use tar to create the archive file. To create the archive, first ensure that your src directory contains only the source code and makefiles needed to compile your project. Any compiled programs, object files, or other additional files should be moved or removed.

Next, move to the directory containing src and execute this command:

```
pa1$ tar -vzcf pa1.tar src
```

tar will create a file pal.tar that contains all files in the directory src. This file can now be submitted through Sakai.

To verify that the archive contains the necessary files, you can print a list of the files contained in the archive with this command:

```
pa1$ tar -tf pa1.tar
```

You should also use the auto-grader to confirm that your archive is correctly structured.

```
pa1$ ./grader.py -a pa1.tar
```

3.5 Using the auto-grader

We have provided a tool for checking the correctness of your project. The auto-grader will compile your programs and execute them several times with different arguments, comparing the results against the expected results.

Setup The auto-grader is distributed as an archive file pal-grader.tar. To unpack the archive, move the archive to a directory and use this command:

```
pa1$ tar -xf pa1-grader.tar
```

This will create a directory pa1 containing the auto-grader itself, grader.py, a library autograde.py, and a directory of test cases data.

Do not modify any of the files provided by the auto-grader. Doing so may prevent the auto-grader from correctly assessing your program.

You may create your src directory inside pa1. If you prefer to create src outside the pa1 directory, you will need to provide a path to grader.py when invoking the auto-grader (see below).

Usage While in the same directory as grader.py and src, use this command:

```
pa1$ ./grader.py
```

The auto-grader will compile and execute the programs in the directory src, assuming src has the structure described in section 3.2.

By default, the auto-grader will attempt to grade all programs. You may also provide one or more specific programs to grade. For example, to grade only sorta:

```
pa1$ ./grader.py sorta
```

To stop the auto-grader after the first failed test case, use the --stop or -1 option.

To obtain usage information, use the -h option.

Program output By default, the auto-grader will not print the output from your programs, except for lines that are incorrect. To see all program output for unsuccessful tests, use the --verbose or -v option:

```
pa1$ ./grader.py -v
```

To see program output for all tests, use -vv. To see no program output, use --quiet or -q.

Checking your archive We recommend that you use the auto-grader to check an archive before submitting. To do this, use the --archive or -a option with the archive file name. For example,

```
pa1$ ./grader.py -a pa1.tar
```

This will unpack the archive into a temporary directory, grade the programs, and then delete the temporary directory.

Specifying source directory If your src directory is not located in the same directory as grader.py, you may specify it using the --src or -s option. For example,

pa1\$./grader.py -s ../path/to/src

Refreshing the build directory In the unlikely event that your build directory has become corrupt or otherwise unusable, you can simply delete it using rm -r build. Alternatively, the --fresh or -f option will delete and recreate the build directory before testing.